



Improving Automatic Fetal Biometry Measurement with Swoosh Activation Function

Shijia Zhou^{1,2}(✉), Euijoon Ahn³, Hao Wang¹, Ann Quinton⁴,
Narelle Kennedy⁵, Pradeeba Sridar⁶, Ralph Nanan^{1,2,3,4,5,6},
and Jinman Kim¹

¹ University of Sydney, School of Computer Science, Sydney, NSW 2000, Australia
szho6430@uni.sydney.edu.au

² James Cook University, Smithfield, QLD 4878, Australia

³ Central Queensland University, Sydney, NSW 2000, Australia

⁴ Liverpool Hospital, Liverpool, NSW 2170, Australia

⁵ Indian Institute of Technology, Delhi, New Delhi, Delhi 110016, India

⁶ University of Sydney, Sydney Medical School Nepean, Kingswood, NSW 2747,
Australia

Abstract. The measurement of fetal thalamus diameter (FTD) and fetal head circumference (FHC) are crucial in identifying abnormal fetal thalamus development as it may lead to certain neuropsychiatric disorders in later life. However, manual measurements from 2D-US images are laborious, prone to high inter-observer variability, and complicated by the high signal-to-noise ratio nature of the images. Deep learning-based landmark detection approaches have shown promise in measuring biometrics from US images, but the current state-of-the-art (SOTA) algorithm, BiometryNet, is inadequate for FTD and FHC measurement due to its inability to account for the fuzzy edges of these structures and the complex shape of the FTD structure. To address these inadequacies, we propose a novel Swoosh Activation Function (SAF) designed to enhance the regularization of heatmaps produced by landmark detection algorithms. Our SAF serves as a regularization term to enforce an optimum mean squared error (MSE) level between predicted heatmaps, reducing the dispersiveness of hotspots in predicted heatmaps. Our experimental results demonstrate that SAF significantly improves the measurement performances of FTD and FHC with higher intraclass correlation coefficient scores in FTD and lower mean difference scores in FHC measurement than those of the current SOTA algorithm BiometryNet. Moreover, our proposed SAF is highly generalizable and architecture-agnostic. The SAF's coefficients can be configured for different tasks, making it highly customizable. Our study demonstrates that the SAF activation function is a novel method that can improve measurement accuracy in fetal

This research was supported by Australian Research Council (ARC) DP200103748.

Supplementary Information The online version contains supplementary material available at https://doi.org/10.1007/978-3-031-43990-2_27.

biometry landmark detection. This improvement has the potential to contribute to better fetal monitoring and improved neonatal outcomes.

Keywords: 2-dimensional ultrasound · automatic measurement algorithm · fetal thalamus · activation function

1 Introduction

The thalamus is a critical brain region that relays and modulates information between different parts of the cerebral cortex, and plays a vital role in signal transmission and processing, including pain recognition and reaction [13]. Abnormal fetal thalamus development, which can disrupt serotonin receptor development, may contribute to the development of later neuropsychiatric disorders [12]. Thalamus development is influenced by maternal factors such as diabetes that suppresses thalamus development after 2 weeks of gestation [16]. To further investigate the relationship between maternal factors and fetal thalamus growth, a dataset containing maternal factors, fetal thalamus diameter (FTD), and fetal head circumference measurements (FHC) is needed. FHC measurement is necessary to normalize FTD against gestational age. However, measuring FTD and FHC manually from 2D-US scans is laborious, prone to high inter-observer variability, and complicated by 2D-US images' high signal-to-noise ratio (SNR) nature due to ultrasound wave's lack of penetration power compared to ionizing radiations such as the x-ray, acoustic shadows cast by highly echogenic objects, and unique noise characteristics due to ultrasound reverberation [3, 4, 9].

In recent years, landmark-based detection approaches based on deep learning have been employed to measure biometrics from fetal US images. They were used to detect measurement key points for brain structures in the fetal brains, and bony structures such as length of the femur or dimensions in the pelvic floors [2, 7, 14]. In these studies, the distance between a pair of landmarks represents the biometry being measured. The current state-of-the-art (SOTA) landmark detection algorithm for measuring biometry in fetal 2D-US images is the BiometryNet proposed by Avisdris et al., which was developed to detect landmarks for measuring the FHC and fetal femur length [2]. BiometryNet is based on High-Resolution Net (HRNet) and it has shown great performances in measuring dimensions of fetal skull and femur bone, outperforming other landmark-based methods [10]. BiometryNet uses dynamic orientation determination (DOD) to enforce a consistent orientation between detected landmarks', i.e. the first landmark is always the left/top measurement key point, and the second landmark is always the right/bottom key point [2].

However, BiometryNet cannot be directly used to measure FTD and FHC due to two specific difficulties. First being that the "guitar-shaped" structure (GsS) by Sridar et al. to measure FTD has similar echogenicity to surrounding brain tissues, resulting in fuzzy boundaries, especially around the wing-tips where measurement landmarks are located [9] (Fig. 1). This difficulty is also observed in 2D-US images of fetal skulls when they appear broken due to unfused bones




Fig. 1. Left: 2D-US image of fetal head from the HC18 dataset [11], note part of the skull is not mineralized and has similar echogenicity to the adjacent uterine tissue. Middle: 2D-US image of a fetal femur [1]. Right: 2D-US image of a GsS for measuring FTD, note the gaps in the skull due to unfused bones.

and acoustic shadows cast by the skull bones themselves (Fig. 1). Second being that the shape of the GsS resembles the silhouette of a guitar and more complex than that of a skull or femur bone (Fig. 2) [9]. These two difficulties causes uncertainties in the localization of measurement landmarks of FTD, resulting in inaccurate measurement of FTD.

To address the above difficulties, we present a novel Swoosh Activation Function (SAF) designed to enhance the regularization of heatmaps produced by landmark detection algorithms. The SAF takes its name from the NikeTM swoosh logo, which resembles its shape (Supplementary Fig. 1). SAF can serve as a regularization term to enforce an optimum mean squared error (MSE) level between a pair of predicted heatmaps (pHs). By doing so, the landmark detection algorithm is compelled to highlight different areas. Additionally, SAF can enforce an optimum MSE between individual pH and a zero matrix to reduce hotspot dispersiveness. Moreover, because SAF does not grow exponentially when the input MSE is higher than the optimum MSE level, it does not hinder algorithm's learning. Consequently, we hypothesize that SAF can enhance landmark detection accuracy and overcome uncertainties arising from the fuzzy edges of GsS by promoting hotspot concentration in pHs .

2 Method

2.1 Swoosh Activation Function

SAF is introduced to optimize pHs by enforcing an optimum MSE between a pair of pHs and a secondary optimum MSE between a predicted heatmap (pH) and a zero matrix (O). We determine the optimum MSE by computing the MSE between a pair of ground truth heatmaps (gHs). Each ground truth heatmap (gH) represents a measurement landmark by a smaller matrix drawn from a Gaussian distribution that is centered at the landmark coordinates with the peak assigned at the value of 1 (Fig. 3.A). Since we determine that gHs represent optimum heatmaps, the MSE between a pair of pHs should approximate the MSE between a pair of gHs . In addition, the secondary optimum MSE between

a pH and a zero matrix is half of the MSE between a pair of gHs , since only one Gaussian distribution is being compared. We demonstrate in Fig. 3.B and C how deviations from this optimum MSE value can lead to incorrect and noisy heatmaps.




Fig. 2. A: a 2D-US fetal brain image that has the GsS (outlined with white curly lines) annotated with manually created bounding box (white box); B: input image of the GsS constrained by the bounding box. The red dots represent the ground truth landmarks of FTD; C: the first heatmap of one of the FTD landmarks with the hottest (red) spot representing the landmark; D: the second heatmap of the other FTD landmark with the hottest (red) spot representing the landmark (Color figure online)

To enforce this pre-determined optimum MSE, we defined SAF as:

$$f(x > 0) = \left(ax + \frac{1}{bx}\right)^c - Min \quad (1)$$

In Eq. 1, Min represents a function that ensures SAF minimizes to 0, and it is defined as $Min = f\left(\sqrt{\frac{1}{ab}}\right)$. The coefficient a determines the slope of SAF around the minimum point in Quadrant 1 of the Cartesian coordinate system (Supplementary Fig. 1). The slope of SAF determines its regularization strength. Coefficient b determines the x -axis coordinate of the minimum point where the x -coordinate of the minimum point correspond to the optimum MSE. Coefficient b is deduced by Supplementary Eq. 1 and the coefficient c is deduced by Supplementary Eq. 2.

2.2 SAF Regularization

The loss function consists of the MSE between pH and gH , and three additional SAF regularization terms. We chose SAF as the activation function to control these regularization terms because SAF's output grows exponentially on either side of the minimum point (Supplementary Fig. 1) where the x -axis represents input values of $MSE(pH_1, pH_2)$, $MSE(pH_1, O)$, and $MSE(pH_2, O)$, while y -axis represents the output values of SAF which minimize to 0 when input values approximate the predetermined optimum MSE. SAF locks the input MSE to the pre-determined optimum value because deviation from this value would result in

a fast increase in the gradient of SAF. The first SAF term regularizes the MSE between a pair of pHs . The next two SAF terms regularize the MSE between each pH and a zero matrix. The equation for the entire loss function is:

$$L(pH_1, pH_2, gH_1, gH_2, O) = MSE(pH_1, gH_1) + MSE(pH_2, gH_2) + \\ SAF(MSE(pH_1, pH_2)) + SAF(MSE(pH_1, O)) + \\ SAF(MSE(pH_2, O)) \quad (2)$$

where pH_1 = the first predicted heatmap; pH_2 = the second predicted heatmap; gH_1 = the first ground truth heatmap; gH_2 = the second ground truth heatmap; O = a zero matrix where the elements are zero.




Fig. 3. A (left column): a pair of ground truth heatmaps with optimum MSE = 0.0061. B (middle column): a pair of predicted heatmaps with low MSE = 0.00003. C (right column): a pair of predicted heatmaps with high MSE = 17.7487

2.3 Datasets

FTD Dataset: The dataset used in this study consists of 1111 2D-US images acquired during the second trimester of pregnancies and confirmed by board-certified ultrasonographers to be suitable for measuring FTD [8]. No additional ethics approval was required. Spatial constraints were provided by manually added bounding boxes around the GsS, verified by the same ultrasonographers. Pycocotools generated two gHs for each pair of measurement landmarks, with hotspots representing landmarks [6]. 5-fold cross-validation was performed. During training, 100 training samples were randomly held out as the validation set.

The intraclass correlation coefficient (ICC) score was computed using IBMTM SPSSTM version 28, with the ICC configuration being Two-Way Random and Absolute Agreement [5].

HC18 Dataset: HC18 dataset is available on the Grand Challenge website [11]. We utilized least squared fitting of an ellipse to determine the center, width, height, and angle of rotation of the elliptical ground truth mask. We used trigonometry to determine the coordinates of the landmarks for the major and minor axes of the ellipse. Spatial constraints in the form of bounding boxes were built by the ground truth mask of head circumference published by the dataset author. Predicted landmark points were used for the major and minor axes to calculate the ellipse width and height for testing. The predicted major and minor axes were assumed to be perpendicular, and their point of intersection was used as the center of the ellipse. Finally, we uploaded the results to the Grand Challenge leaderboard and used the mean difference between predicted FHC and ground truth as the performance metric.

2.4 Experimental Setup

Training Epochs and Learning Rate. We conducted our experiments using PyTorch version 1.12 on two NVIDIA GTX-1080Ti graphical processing units, each with 11 GB of video memory. For landmark detection training, we used the same learning rate configuration as Avidris et al. [2] to train both BiometryNet with/without SAF. Specifically, we set the initial learning rate to 10^{-5} and reduced it using a multi-step learning rate scheduler that scaled the learning rate by a factor of 0.2 at epoch numbers 10, 40, 90, and 150. We trained all BiometryNet models for a total of 200 epochs. For EfficientNet with/without SAF, we set the initial learning rate to 10^{-5} and reduced it using a multi-step learning rate scheduler that scaled the learning rate by a factor of 0.1 at epoch number 300 and 350. We trained both models for 400 epochs.

Pre-processing. Our pre-processing pipeline included random rotation of $\pm 180^\circ$, random re-scaling of $\pm 5\%$, resizing to 384×384 pixels without preserving aspect ratio, and normalization using ImageNet-derived mean = (0.485, 0.456, 0.406) and standard deviation = (0.229, 0.224, 0.225) for each color channel.

SAF Configuration. Given our dataset configuration where each biometry landmark was represented by a 19×19 matrix with values derived from a Gaussian distribution with the center point's value peaked at 1, the MSE between a pair of *gHs* was 0.0061. This configuration followed the standard implementation used in human pose estimation landmark detection [15]. The secondary optimum MSE between a *pH* and a zero matrix was halved at 0.00305. We also predetermined the value of *Min* to be 0.001 to prevent SAF from overpowering the MSE loss between *gHs* and *pHs*. We experimented with different values

of the coefficient a (1, 4, and 8) to evaluate coefficient a 's influence on landmark detection accuracy. We then determined the value of coefficient b using Supplementary Eq. 1, and the value of coefficient c using Supplementary Eq. 2.

We conditionally activated SAF when the average of $MSE(pH_1, gH_1)$ and $MSE(pH_2, gH_2)$ was less than 0.0009 because SAF is not bounded and early activation would hinder algorithm learning. The proposed SAF algorithm was evaluated using six model configurations, including Vanilla BiometryNet, BiometryNet with SAF with coefficient a values of 1, 4, and 8, an EfficientNet, and the EfficientNet with SAF configured with coefficient a value of 4. The model configurations were trained and tested to verify the usefulness of the proposed SAF using both FTD and HC18 datasets.

3 Results

The results of FTD dataset show that BiometryNet with SAF_a1 (BiometryNet_SAF_a1) achieved the highest ICC score at 0.737, surpassing the performance of the vanilla BiometryNet, which scored 0.684. Moreover, BiometryNet with SAF_a4 (BiometryNet_SAF_a4) and BiometryNet with SAF_a8 (BiometryNet_SAF_a8) also demonstrated superior ICC scores for FTD measurement, albeit to a lesser degree than BiometryNet_SAF_a1. The impact of SAF on performance was further observed in the modified EfficientNet, where EfficientNet_SAF_a4 achieved a higher ICC score of 0.725, compared to the modified




Fig. 4. Row A: landmarks predicted and heatmaps produced by BiometryNet. Left, input image overlaid with predicted landmarks (red spots). Middle, the first predicted heatmap, there are hotspots present near both upper and lower wing-tips. Right, the second predicted heatmap. Row B: landmarks predicted and heatmaps produced by BiometryNet_SAF_a1. Left input image overlaid with predicted landmarks (red spots). Middle, the first predicted heatmap. Right, the second predicted heatmap. (Color figure online)

EfficientNet without SAF, which only scored 0.688. We also observed that SAF reduced the similarities between a pair of pHs and the dispersiveness of hotspot in the pH . We display such heatmaps produced by BiometryNet in Fig .4.A and BiometryNet_SAF_a1 in Fig .4.B.

For the HC18 dataset, BiometryNet_SAF_a8 demonstrated the lowest measurement mean difference from the ground truth at $3.86 \text{ mm} \pm 7.74 \text{ mm}$. Additionally, all configurations outperformed the vanilla BiometryNet. The impact on FHC measurement was further observed in the modified EfficientNet, where EfficientNet_SAF_a4 achieved a lower measurement mean difference at $4.87 \text{ mm} \pm 5.79 \text{ mm}$ compared to EfficientNet at $32.76 \text{ mm} \pm 21.01 \text{ mm}$. The FTD dataset ICC scores and HC18 dataset mean differences for each algorithm are presented in Table 1.

Table 1. ICC scores of the FTD dataset, and mean measurement differences \pm confidence interval (CI) of the HC18 dataset achieved by all the algorithms.

Network Name	FTD ICC	HC18 Mean Difference \pm CI (mm)
Vanilla BiometryNet	0.684	4.56 ± 7.41
BiometryNet_SAF_a1	0.737	4.02 ± 6.70
BiometryNet_SAF_a4	0.724	4.03 ± 7.73
BiometryNet_SAF_a8	0.719	3.86 ± 7.74
EfficientNet	0.688	32.76 ± 21.01
EfficientNet_SAF_a4	0.725	4.87 ± 5.79

4 Discussion

The main findings of this study are as follows: (1) SAF improved the measurement accuracy of algorithms in both FTD and FHC measurement tasks; (2) SAF regularization is architecture-agnostic, as it improved the measurement accuracy of both BiometryNet and EfficientNet compared to their vanilla forms that do not use SAF; and (3) the optimum configuration of SAF coefficients is task-dependent. For FTD measurement, the most optimum configuration was to use $a = 1$, while for FHC measurement was with coefficient $a = 8$.

The results demonstrate the performance improvement brought about by SAF regularization, effectively improving the accuracy of landmark detection of fetal biometries in 2D-US images. SAF regularization forces a pair of heatmaps to highlight different areas and reduce the dispersiveness of hotspots in pHs , which results in improved fetal biometry landmark detection accuracy. This is supported by the comparison of heatmaps produced by BiometryNet and BiometryNet_SAF_a1 displayed in Fig. 4. Moreover, SAF regularization is simple to implement and easy to configure, requiring no modification to the network architecture. Our results also suggest that SAF is highly generalizable

as it is architecture-agnostic, improving the performance of both BiometryNet and EfficientNet. SAF is also highly configurable for different tasks via different coefficient configurations. Furthermore, we also suggest that SAF is generalizable to other imaging modalities that also require pair-wise landmark detection. For example, detecting mitral and aortic valves in 2D-US images of heart and detecting cranial sutures in CT images of skull.

As part of our future study, we will explore the effectiveness of SAF regularization in other fetal landmark detection tasks, especially those that suffer from similar issues with fuzzy edges and uncertain landmark locations. Additionally, the optimum configuration of SAF coefficients may vary depending on the specific dataset or imaging modality used.

5 Conclusion

Our study demonstrated the effectiveness of SAF as a novel activation function for regularizing heatmaps generated by fetal biometry landmark detection algorithms, resulting in improved measurement accuracy. SAF outperformed the previous state-of-the-art algorithm, BiometryNet, in both FTD and FHC measurement tasks. Importantly, our results showed that SAF is architecture-agnostic and highly configurable for different tasks through its coefficients, making it a generalizable solution for a wide range of landmark detection problems.

References

1. Apostolos Kolitsidakis: How to measure the femur length (2021)
2. Avidris, N., et al.: Biometrynet: landmark-based fetal biometry estimation from standard ultrasound planes. In: Wang, L., Dou, Q., Fletcher, P.T., Speidel, S., Li, S. (eds.) MICCAI 2022, vol. 13434, pp. 279–289. Springer, Cham (2022). https://doi.org/10.1007/978-3-031-16440-8_27
3. Bethune, M., Alibrahim, E., Davies, B., Yong, E.: A pictorial guide for the second trimester ultrasound. *Australas. J. Ultrasound Med.* **16**(3), 98–113 (2013). <https://doi.org/10.1002/j.2205-0140.2013.tb00106.x>
4. Brickson, L.L., Hyun, D., Jakovljevic, M., Dahl, J.J.: Reverberation noise suppression in ultrasound channel signals using a 3D fully convolutional neural network. *IEEE Trans. Med. Imaging* **40**(4), 1184–1195 (2021). <https://doi.org/10.1109/TMI.2021.3049307>
5. IBM Corp: IBM SPSS Statistics for Windows (2020)
6. Lin, T.-Y., et al.: Microsoft COCO: common objects in context. In: Fleet, D., Pajdla, T., Schiele, B., Tuytelaars, T. (eds.) ECCV 2014. LNCS, vol. 8693, pp. 740–755. Springer, Cham (2014). https://doi.org/10.1007/978-3-319-10602-1_48
7. Shankar, H., et al.: Leveraging clinically relevant biometric constraints to supervise a deep learning model for the accurate caliper placement to obtain sonographic measurements of the fetal brain. In: *Proceedings - International Symposium on Biomedical Imaging*, vol. 2022-March. IEEE Computer Society (2022). <https://doi.org/10.1109/ISBI52829.2022.9761493>

8. Sridar, P., Kennedy, N.J., Quinton, A.E., Robledo, K., Kim, J., Nanan, R.: Normative ultrasound data of the fetal transverse thalamic diameter derived from 18 to 22 weeks of gestation in routine second-trimester morphology examinations. *Australas. J. Ultrasound Med.* **23**(1), 59–65 (2020). <https://doi.org/10.1002/ajum.12196>
9. Sridar, P., et al.: Automatic measurement of thalamic diameter in 2-D fetal ultrasound brain images using shape prior constrained regularized level sets. *IEEE J. Biomed. Health Inform.* **21**(4), 1069–1078 (2017). <https://doi.org/10.1109/JBHI.2016.2582175>
10. Sun, K., Xiao, B., Liu, D., Wang, J.: Deep high-resolution representation learning for human pose estimation. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 5693–5703 (2019)
11. van den Heuvel, T.L., de Bruijn, D., de Korte, C.L., Ginneken, B.V.: Automated measurement of fetal head circumference using 2D ultrasound images (2018)
12. Wai, M.S., Lorke, D.E., Kwong, W.H., Zhang, L., Yew, D.T.: Profiles of serotonin receptors in the developing human thalamus. *Psychiatry Res.* **185**(1–2), 238–242 (2011). <https://doi.org/10.1016/j.psychres.2010.05.003>
13. Waxman, S.G., Waxman, S.G.: Chapter 9: Diencephalon: Thalamus and Hypothalamus. In: *Clinical Neuroanatomy*, chap. 9. McGraw-Hill Education, New York, 29th edn. (2020)
14. Xia, W., et al.: Automatic plane of minimal hiatal dimensions extraction from 3D female pelvic floor ultrasound. *IEEE Trans. Med. Imaging* **41**(12), 3873–3883 (2022). <https://doi.org/10.1109/TMI.2022.3199968>
15. Xiao, B., Wu, H., Wei, Y.: Simple baselines for human pose estimation and tracking. In: *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 466–481 (2018)
16. You, L., Deng, Y., Li, D., Lin, Y., Wang, Y.: GLP-1 rescued gestational diabetes mellitus-induced suppression of fetal thalamus development. *J. Biochem. Mol. Toxicol.* (2022). <https://doi.org/10.1002/jbt.23258>